

PAPER TALK · SEPTEMBER 2026

Talking at the Audience

Listener-directed address in political podcasts, and the institutional cynicism of the people who listen

Danielle Yang with R. Michael Alvarez

Linde Center for Science, Society, and Policy · Caltech

A major political medium, mostly unmeasured

- Podcasts are now a primary political news source for millions of Americans, and the largest political shows out-draw cable news.
- Almost all of what we know about partisan media comes from television and print: incivility, outrage, strategy framing.
- Podcasts are different in form: long, intimate, one voice speaking directly to a listener for hours a week.
- The common assumption is that the right-leaning shows are angrier. Nobody had measured it at scale.
- And nobody had linked what hosts say to who is actually listening and what those listeners believe.

TWO GAPS THIS PAPER FILLS

1 Content. A transcribed, validated corpus of the political podcast field, with register measured on it.

2 Audience. The corpus linked to two national probability surveys, so shows can be matched to their own listeners.

Three questions

RQ1 How do left- and right-leaning political podcasts differ in how hosts talk?

Hostility? Moral language? Or something about the relationship the host sets up with the listener?

RQ2 Who listens to political podcasts, and how do they differ from other news consumers?

Institutional trust, democratic norms, partisanship, participation — measured in national surveys.

RQ3 Does a show's mode of address predict who its audience is and what they believe?

The show-level link: register measured on transcripts, outcomes measured on the show's own listeners.

Three literatures, one prediction each

Literature	Anchor	What it predicts here	What we find
Media malaise / spiral of cynicism	Robinson 1976; Cappella & Jamieson 1997; Ladd 2012	News diets outside institutional journalism go with lower institutional trust	Supported (RQ2)
Incivility and outrage	Mutz & Reeves 2005; Sobieraj & Berry 2011	Right-leaning shows are more hostile; hostility drives distrust	Not supported for the partisan contrast
Para-social interaction	Horton & Wohl 1956	Direct personal address builds a one-sided relationship with the performer	Consistent: address, not hostility, is the partisan difference
Register variation	Biber 1988	Involved vs. informational style is a stable property of a speaker	Address is a stable show trait; reaches less-educated listeners

Populist communication (Rooduijn & Pauwels 2011; Fawzi 2019) enters as a rival content feature, carried as a covariate rather than ruled out.

The corpus

204

political podcasts from the Apple Podcasts US Politics chart, ranks 1-250, frozen 13 July 2026

29,421

episodes transcribed (faster-whisper large-v3-turbo, two GB10 nodes, ~100x real time)

1.38 M

750-character passages, the unit the ideology instrument was validated on

2018-

2026
with a census of the pre-2018 tail; episodes drawn per show per quarter so show-time is aligned

- Inclusion filter applied in fixed order: 250 charted → 219 (non-US-politics, under 10 hours of audio, no usable feed) → 204 on the topic-arm sample.
- Scope: political podcasts that attained and held Politics-chart presence. Mid-tail by construction; the largest political shows (Shapiro, The Daily, Tucker) are filed under other categories and enter as an out-of-frame test.
- Ideology labelled at passage level by a two-stage LLM classifier (152,152 passages), aggregated to show level; host ideology anchored externally with DIME campaign-contribution scores (112 shows).

Two national surveys, linked to the corpus

Survey	Sample	Exposure	Outcomes
Kettering Foundation / Gallup, Democracy for All, Year 1 (Jul–Aug 2025)	20,338; 45% Gallup Panel (probability), 55% opt-in; weighted	Open-ended top-three news sources: 51,082 verbatims. Matched to shows by exact title (385 respondents), plus a re-scan for misspellings (150), plus six out-of-frame political podcasts (290)	8-item institutional cynicism ($\alpha = .84$), election distrust, democratic-norms battery, participation
Pew American Trends Panel, Waves 150 / 155 / 165 (Jul 2024 – Mar 2025)	10,658 in W150; ~9,000 linked to each later wave; probability panel	Frequency of podcast news; podcasts as main election-news source; seven items on relationships with news influencers	Trust in national news, local news, social media, and friends/family — measured 2–8 months after exposure

Kettering lets us match a show's transcripts to that show's own listeners. Pew lets us put exposure before outcome.

Measuring how a host talks to the listener

- **Listener-directed address** = mean of three show-level z-scores, each an n-weighted rate per 10,000 words:
 - syntactic imperatives — a bare-form verb heading a clause with no subject or auxiliary (spaCy dependency parse)
 - dictionary imperatives — look, listen, think about it, remember, understand ...
 - second-person pronouns — you, your, yourself
- **Hostility** from Sobieraj & Berry's outrage categories plus a profanity list; moral language from eMFD and MFD 2.0; Biber dimensions via the Multi-Feature Tagger of English.
- All instruments published before this study. No word list was written for it.
- Register scored on every passage (ads flagged, not removed); topic model run on ad-excised text.

WHAT THE MEASURE PICKS UP

“You need to understand what they’re doing here. Look at what happened last week. Don’t let them tell you otherwise.”

Deontic second person (you need to, you should) is the strongest single component against host ideology ($r = +0.44$). The filler ‘you know’ is null. Biber’s persuasion dimension is uncorrelated ($r = -0.12$). The second-person subject carries the signal, not the modality.

Ideology labels validated four ways

Standard	Basis	n	Agreement
DIME cfscore	hosts' campaign contributions	112 shows	$r = +0.68$, $\kappa = 0.66$
Editorial ratings	human-assigned partisan leaning	25	$\kappa = 0.75$
Expert ratings	independent research coding	26	$\kappa = 0.77$
Audience party ID	party of the show's own Kettering listeners	22 shows / 494 listeners	$r = +0.91$ $[0.77, 0.98]$

21 of 22

shows whose label predicts the majority party of their own listeners

Left-labelled shows average 76% Democratic audiences; right-labelled 91% Republican.


Text, money and listeners agree. This validates show labels; a 200-passage two-coder check of passage labels precedes submission.

Inference built for few, unequal clusters

- **Between-show register differences:** show-level permutation tests, Benjamini-Hochberg across measure families, format and ad controls, a discriminant null (filler 'you know'), and equivalence (TOST) tests for the nulls.
- **Population survey estimates:** weighted OLS with robust SEs, full control ladder (party, attention, age, education, gender, race, income, social-media use, platforms, urbanicity, mode, number of sources), within-party and within-education re-estimation.
- **Show-linked audience estimates:** exposure is assigned by show and a few shows hold most listeners, so conventional clustered SEs are anti-conservative. Every estimate reports CR2 (Bell-McCaffrey) with Satterthwaite df as the primary p and a wild cluster bootstrap alongside; the df (≈ 9 on the full frame) is the honest cluster count.
- **Pre-registration:** the out-of-frame test was frozen and hashed before any audio was fetched — hypothesis, alias table, outcome, model, inference, failure criterion.

RESULTS

Finding 1 — Left and right differ in mode of address, not hostility



Pearson r between each show-level register measure and the host's DIME contribution score.
Grey bars: not significant.

$d = 0.78$

right vs left shows on listener-directed address, format-controlled ($p < .0001$)

Insult: $r = 0.00$

tracks ideological intensity ($r = +0.55$), not direction

AUC 0.80 / 0.57

a register classifier separates left from right between shows, barely within them: register is a property of the show

RESULTS

Finding 1, continued — It is how they talk, not what is happening


- **Eight political shocks, 2020-2025:** two general elections, a midterm, the inauguration, the assassination attempt, Biden's withdrawal, January 6, Dobbs.
- Within-show event study: listener-directed address does not shift on either side (all $|b| \leq 0.08$ SD).
- The 2024 election 'shift' is indistinguishable from 60 placebo dates ($p = .57$).
- Across a 75-topic model, within-show topic means spread only 0.04 SD; the most second-person passages are ad reads.
- Hosts do not ramp it up before elections and do not change it by subject. It is a trait of the host.

WHY THIS MATTERS

If address were a reaction to the news cycle, it would be a poor candidate for a stable audience-selection mechanism. Because it is invariant to events and topics, a listener who chooses a show is choosing a mode of address that will be there every week.

RESULTS

Finding 2 — Podcast news diets go with institutional cynicism



Kettering Y1, 8-item cynicism vs mainstream-only (first three) or vs all other news consumers; weighted; full control ladder. All $p < .001$.

-0.10 SD

for each additional mainstream source named ($p = 2e-21$); +0.08 for each additional podcast


+0.14 SD

podcast-only beyond platform-only; equivalence rejected at ± 0.20

$n = 18,257$ · coefficient moves 3% across the control ladder · identical at every education level (interaction $p = .91$)

RESULTS

Finding 2, continued — Replicates in Pew with exposure measured first



Trust, SD per step of podcast-news frequency measured July-Aug 2024 (1 never – 4 often); $n \approx 9,000$ per wave; weighted, controls. First three bars $p < .0001$; friends and family ns.

–0.34 SD


trust in national news, March 2025,
among those who said podcasts were
their main election-news source

The shape is the point:

institutional trust falls, trust in the
informal channel rises, interpersonal
trust is flat. A response-style artifact
cannot produce that pattern. The
stricter exposure gives effects three to
four times larger.

RESULTS

Finding 2, continued — Engaged anti-institutionalism, not disengagement




Democratic-norms battery, SD units, full controls, BH-corrected across all 115 items.
Platform-only shows none of the civil-libertarian pattern (bars at zero).

Where the two non-institutional diets diverge	Podcast-only	Platform-only
Registered to vote	no difference	lower
Tolerance of political violence	no difference	higher
National identity	no difference	weaker
Loneliness	no difference	—

Podcast-only listeners distrust institutions and are more protective of civil liberties, no more partisan, no lonelier, no less registered. Platform-only diets carry weaker national identity, income strain, lower registration and more tolerance of political violence.

RESULTS

Finding 3 — Listener-directed shows reach less formally educated audiences



Party and show lean controlled; CR2 and wild-bootstrap $p \leq .048$ in every sample. Unchanged with age, gender, income, race, urbanicity and attention added (-0.28).

4 of 4

samples replicate, including the out-of-frame shows where nothing else did


Education specifically

income, age and gender are null; holds within left shows (-0.29) and right shows (-0.19); stable to dropping any of the 12 largest shows; show-level $r = -0.57$

First-source namers -0.35 vs later -0.20 : the fingerprint of selection.

RESULTS

Finding 4 — Podcast news use carries the para-social signature



Pew W150, regular consumers of news from influencers (n = 2,012); BH-corrected. Grey: not significant (q = .19, .69).


26% → 39%

share who feel a personal connection to an influencer, from never to often getting news from podcasts; holds within both parties

The relational and opinion items move; the informational items do not.

That is the audience-side signature of direct, personal address.

Finding 5 — Within political podcasts, address goes with election distrust



SD per SD of listener-directed address; full controls (party, attention, age, education, gender, race, income, urbanicity, right-wing platforms, show lean, frame). Election distrust and cynicism are the two associations that survive populist vocabulary and ideological intensity in the same model.

Pooled (56 shows, 792 listeners)	b / SD	CR2 / wild p
Election distrust	+0.267	< .001 / < .001
Institutional cynicism	+0.221	.005 / < .001
Disaffection	+0.159	< .001 / .003
Economic grievance	+0.182	.007 / .005
Exclusivity	+0.089	.066 / .055

Political podcast = at least 15% of passages political (corpus 10th percentile), applied to corpus and out-of-frame shows alike. Election distrust stays +0.20 to +0.29 at every cutoff from 8% to 30% and survives dropping any show; the six out-of-frame shows carry only 1.7 effective df alone.

What was pre-registered, and what happened

Test	Frozen	Outcome
Out-of-frame generalization: address → cynicism among 749 fresh listeners of ten large chart-absent shows	2 Sept 2026, hashed, before any audio was fetched; two amendments filed before any outcome was opened	Did not pass ($b = -0.001$; Rogan-excluded $+0.22$, $p = .08$). Reported as a result.
Within-frame fresh-listener test: 150 corpus-show listeners recovered by re-scan	3 Sept 2026, before outcomes; Majority Report scored blind	Did not pass ($b = -0.06$).
Q28 free-text test ('what does democracy mean to you')	Timestamped before the variable was opened	Did not pass ($p = .125$).

Everything after those tests is labelled exploratory: the political-content frame, the education-composition result, and the comparison with populist vocabulary and ideological intensity were found by searching and are reported with the search disclosed (an outcome-scan family table is in the methods).

The three content features inter-correlate at 0.4–0.5; with about 56 clusters we do not claim to have adjudicated among them beyond finding 5.

Scope and limitations

- **Association, not effect.** People choose their podcasts. The Pew lag gives temporal order, not identification. No dose-response on any attitude; the education gradient sharpens with dose, which is what selection predicts.
- **Scope.** Politics-chart shows plus the six largest political podcasts outside it; population defined by content. The content definition is post hoc and pre-registered for Year 2.
- **Address is format-sensitive.** In interview shows the second person is often addressed to the guest. The partisan contrast survives removing the interrogative component; cross-format levels need diarization, which the archived audio permits.
- **Content features are collinear.** Address, populist vocabulary and ideological intensity co-vary across shows and are jointly associated with audience disaffection; Year 1 ranks them only within the political frame.
- **Validation gaps.** Show labels validated four ways; passage labels not yet human-coded. Ad-excision gold set annotated by an LLM, human re-annotation pending.
- **Small effective cluster counts,** reported honestly through the Satterthwaite df.

TAKEAWAYS

What the paper claims

How the sides differ

Partisan difference in political podcasting lies in how hosts address their audience — second person and instructions — not in how hostile they are. It is a stable trait of the show, validated against donations, experts and listeners' own party.

Who listens

People who get political news from podcasts are about 0.22 SD more cynical about institutions than other news consumers, in two probability panels, one lagged, in both parties and at every education level — and they are civil-libertarian, not disengaged.

What address predicts

Within political podcasts, the hosts who talk at the listener reach less-educated audiences (robust) who distrust elections and institutions (exploratory, pre-registered for Year 2).

What comes next

- **Before submission**

- 200-passage, two-coder human check of ideology labels
- Human re-annotation of the 886-sentence ad gold set
- Disclosure paragraph and family table in the methods
- Freeze corpus counts; binary outcomes in percentage points

- **Kettering Year 2 pre-registration**

- Content-defined frame stated in advance
- Election distrust as the primary outcome for address
- Populist vocabulary and intensity as covariates; interview dates requested for within-show designs

- **Measurement**

- Diarize archived audio for interview shows; rescore address on host turns only
- Band 6–10 corpus extension (in progress) for quarter-level analyses

- **The causal study**

- Survey experiment: the same segment with and without second-person and imperative address, randomly assigned; efficacy and election-trust items after
- The only design that answers whether address does anything to a listener rather than who it attracts

Thank you

Danielle Yang · dsyang@caltech.edu

Linde Center for Science, Society, and Policy, Caltech · with R. Michael Alvarez

Corpus: 204 shows · 29,421 episodes · 1.38M passages · Kettering-Gallup Democracy for All Y1 · Pew ATP W150/155/165

Corpus construction and transcription, in detail

- Frame: Apple Podcasts US Politics chart, ranks 1-250, frozen 13 July 2026 (250 is as deep as Apple serves).
- Filter in fixed order: no feed URL (2), feed unreachable (2), unparseable (1), non-US-politics by blocklist/regex/language/curated list (14), under 10 hours of audio (12) → 219; 15 not in the topic-arm transcript sample → 204.
- Sampling: census of every episode before 2018, then a fixed number of episodes per show per quarter, ranked by a frozen hash so larger draws nest earlier ones.
- Transcription: faster-whisper batched pipeline, large-v3-turbo (CTranslate2, float16), batch 64, English fixed, Silero VAD, no diarization; two NVIDIA GB10 nodes at $\sim 100\times$ real time.
- Register arm: exhaustive 750-character chunks at sentence boundaries, nothing dropped; ads flagged. Topic arm: sentence-level ad/boilerplate excision (5.9% of words; gate $\geq .90$ precision, $\leq 2\%$ content false positives) then ~ 500 -content-word passages.
- Ideology: two-stage gated LLM classifier at passage level; DIME cfscores for hosts via a three-stage matching cascade with logged manual resolution.

Why CR2 and the wild bootstrap

- On the 56-show political-podcast frame, 57% of clusters hold fewer than five listeners and the four largest (MeidasTouch, Shapiro, NBC Nightly News, The Daily) hold 47%.
- Conventional CR1 standard errors are anti-conservative in that regime (Cameron, Gelbach & Miller 2008); a naive p of .0009 on one exclusivity estimate became .17 under the wild bootstrap.
- Every show-clustered estimate therefore reports CR2 (Bell-McCaffrey) with Satterthwaite degrees of freedom (Pustejovsky & Tipton 2018) as the primary p, and a Rademacher wild cluster bootstrap (1,000–2,000 replications) alongside. The two agree on every estimate in the paper.
- A show-level precision-weighted permutation of the exposure is reported as a low-power third check; leave-one-show-out and content-threshold sensitivity accompany every headline.
- Composites are standardized on the full Kettering sample, never on the listener subsample; address is standardized on the 194-show corpus reference, and out-of-frame shows are scored against that same reference.